

HumVI: Image Stretching and Scaling Tests

Phil Marshall*, Cato Sandford, David Hogg, Amit Kapadia

December 20, 2012

Abstract

We investigate some useful input parameters for use when composing color images from the data in the CFHTLS.

1 Introduction

Suppose we have a large set of imaging data from a given sky survey. We would like to make color representations of the images in this set, such that they can be a) compared with each other, to build intuition about the data quality and the appearance of objects in the survey, and b) searched for low surface brightness, color-contrasting features. We use the HUMVI python implementation of the Lupton et al. (2004, hereafter L04) algorithm, with some simplifications and extensions. For simplicity we use a single filter for each RGB channel, choosing the i , r and g bands respectively.

2 Scaling and Stretching the Images

The scaling and stretching of the input images is controlled by three parameters. The three `scales` parameters, one for each channel, are used to multiply the channel images before any other operations are performed. The `scales` account for any difference in units between the images, and also the sensitivity of the detector in that filter, the exposure time used, and so on. We denote this scale parameter by s_X where X is one of the channel identifiers, R, G, B . For convenience, we normalize the scales to have unit mean, since it will often be the case that the images taken in different filters will have the same units and approximately equal pixel values. Crucially, these scales can be chosen to be the same for all images in a survey, allowing different images to be compared with each other; likewise they can be used to account for variations in, for example, exposure time between images.

After scaling, the total intensity image is computed, and used to compute the stretch factor, which

*dr.phil_marshall@gmail.com

is governed by two parameters, Q and α as given by L04:

$$I = rs_R + gs_G + bs_B \quad (1)$$

$$X(I) = \frac{1}{QI} \cdot \operatorname{arcsinh}(\alpha QI) \quad (2)$$

$$R = rs_RX \quad (3)$$

$$G = gs GX \quad (4)$$

$$B = bs BX \quad (5)$$

$$(6)$$

For small values of αQI , $X \approx \alpha$, and constant: at low intensity, each channel image is simply rescaled by α . Low surface brightness features are made more visible by increasing α . At higher values, the $\operatorname{arcsinh}$ function reduces this scale factor, making high intensity regions less saturated. The onset of this behavior occurs when $\alpha QI \approx 1$, or when $I \approx 1/\alpha Q$. In order to make a PNG image, HUMVI works to make three channel images whose values are clipped at zero and one. This choice of 1 as the maximum pixel value allows us to choose Q and α sensibly. For example, suppose we have a set of scaled images with approximately zero mean, unit rms and brightest pixel value 10^4 . If we want a pixel with value 3 times the rms to have normalized value 0.1 in each channel of the final image, then we need $X(9) = 1/30$ ($I \approx 9$ if each channel image pixel value is about 3σ). We'd like this to still be in the linear regime, so we need $9\alpha Q \ll 1$, and also $\alpha \approx 1/30 \approx 0.03$. Combining these two requirements, we find that we need $Q \ll 3$. The algorithm is not very sensitive to the value of Q , as long as it is very much less than $1/\alpha$. Drawing a parallel with television controls, Q behaves like the brightness, while α is like the contrast.

In Figure 1 we show, for a fixed set of scales, the effect of varying α and Q when displaying an image that has approximately unit rms pixel value in each channel. The values $\alpha = 0.03$ and $Q = 1$ provide a good representation of the image.

3 Saturation and Thresholding

The scaled and stretched pixel values of the previous section have to be mapped onto a unit range for encoding in a PNG image. How we deal with pixels that fall outside that range will affect the appearance of the composite.

At low brightness, we have to decide which pixels we want to appear black. Background-subtracted images will have negative pixel values in the “blank” sky regions. One option is to set all pixels with value less than zero to zero. This leads to a large number of black pixels (approximately half of them!) and a strong impression of dark sky. If we want to retain the information that those negative pixels contain (about the noise level), we can add an offset δ to each scaled and stretched channel image, such that pixels with value $-\delta$ appear black in the final composite. This has the effect of making the blank regions of the image appear dark gray instead of black. Choosing δ to be negative has the opposite effect, making the sky look blacker than black... Figure 2 shows the effect of various offset values on the appearance of the test image from Figure 1. We find that actually an offset of zero is a good compromise between “seeing the noise” and achieving a nice dark background against which low surface brightness features can be seen.

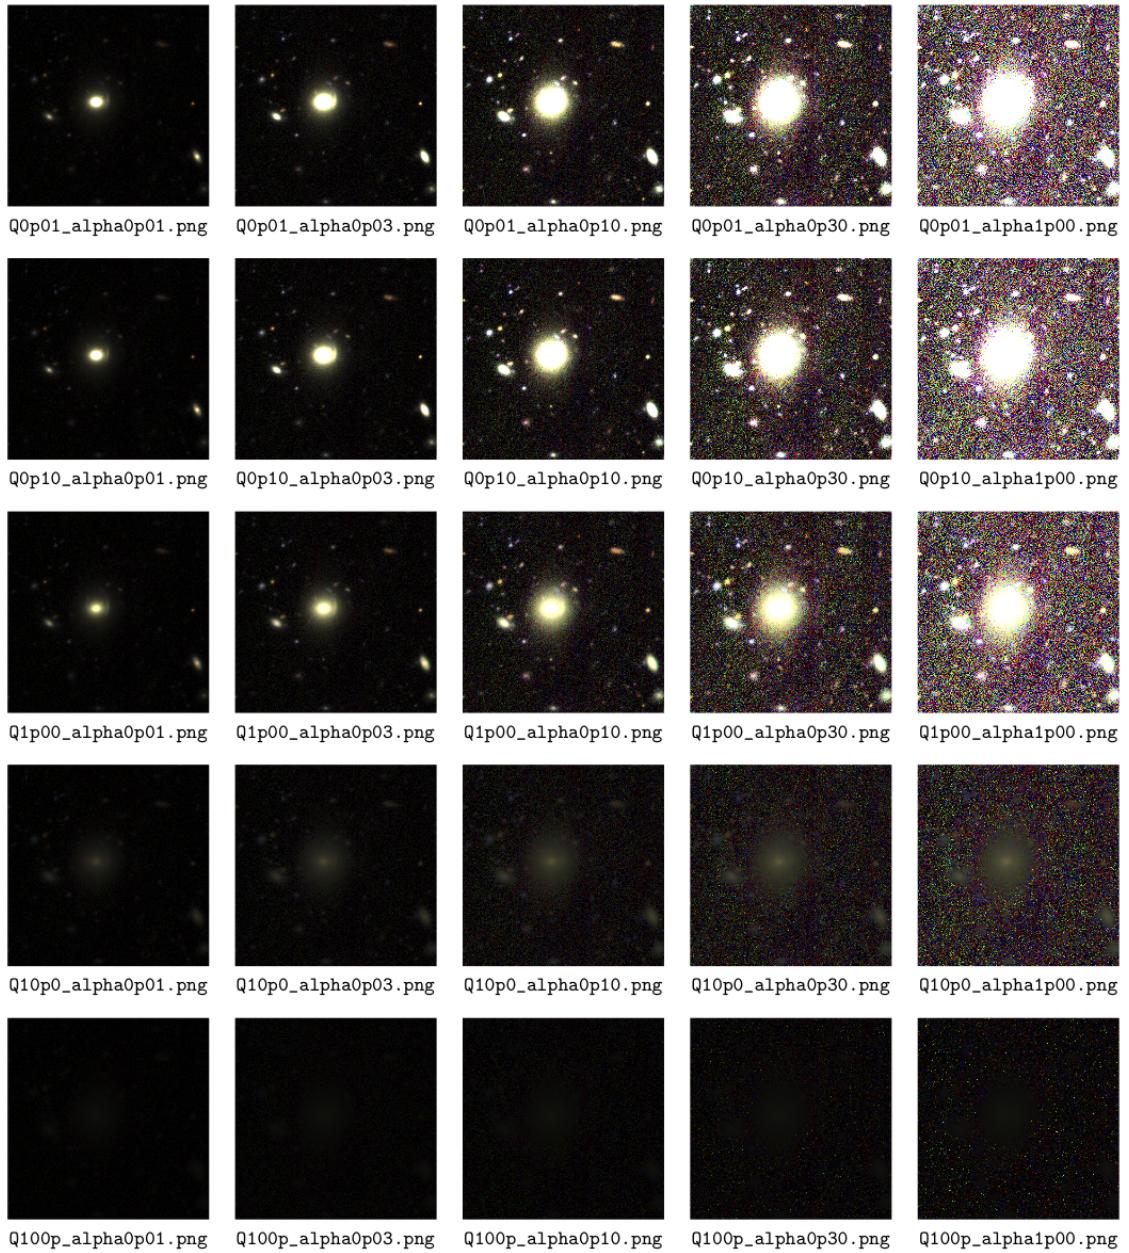


Figure 1: The effect of the non-linearity parameters Q and α on an example image from the CFHTLS survey. Left to right, α increases through the set $\{0.01, 0.03, 0.1, 0.3, 1.0\}$. Top to bottom, Q increases through the set $\{0.01, 0.1, 1.0, 10, 100\}$.

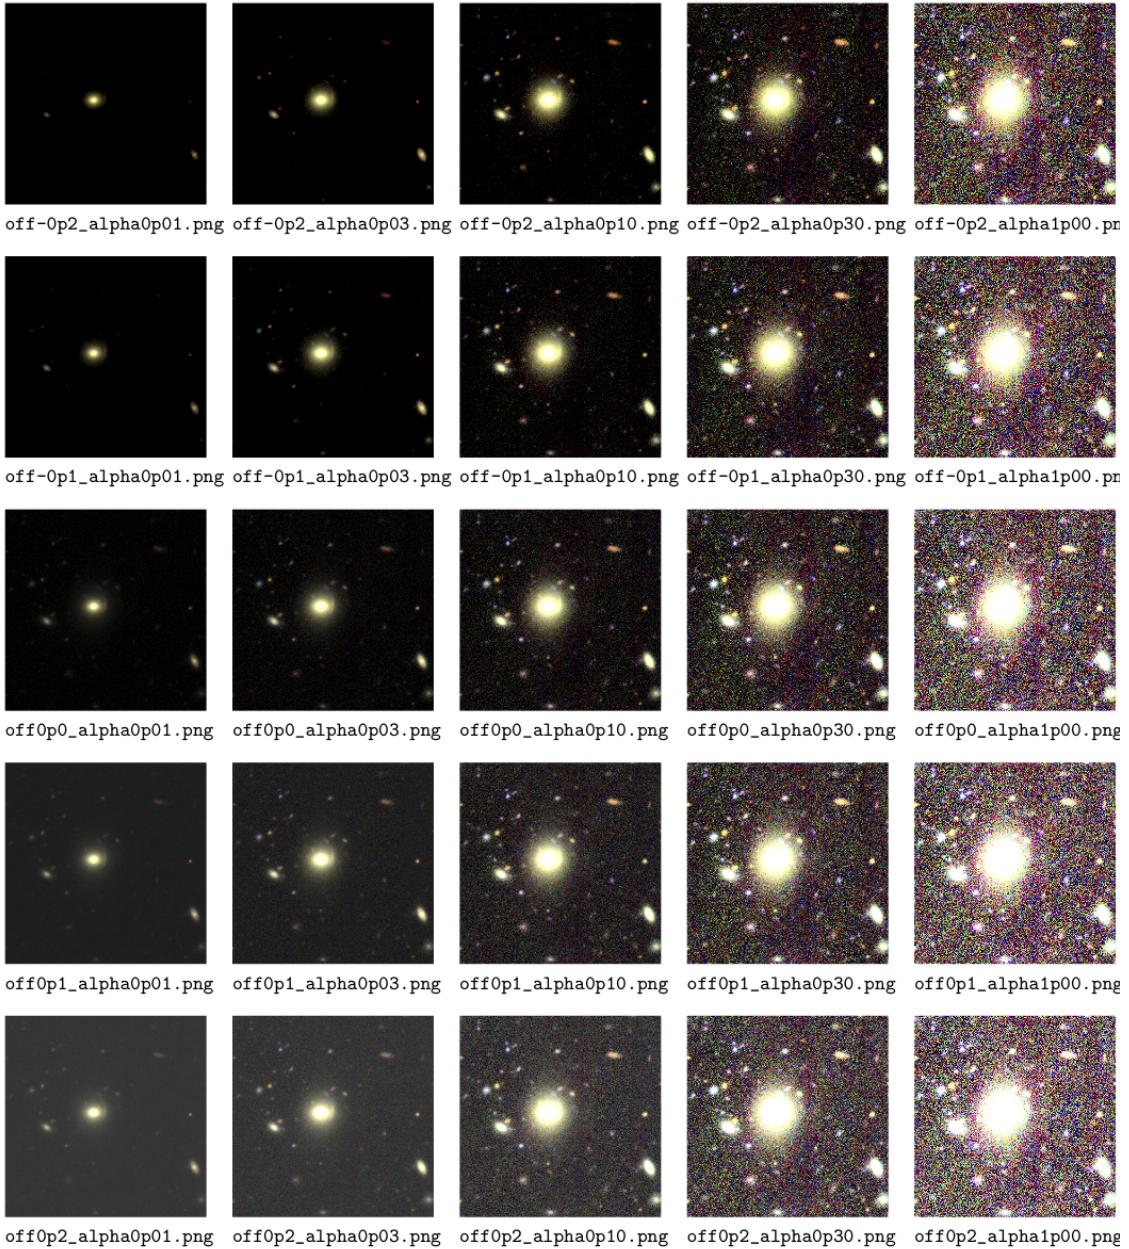


Figure 2: The effect of the offset parameter δ and contrast α on an example image from the CFHTLS survey. Left to right, α increases through the set $\{0.01, 0.03, 0.1, 0.3, 1.0\}$. Top to bottom, δ increases through the set $\{-0.2, -0.1, 0.0, 0.1, 0.2\}$.

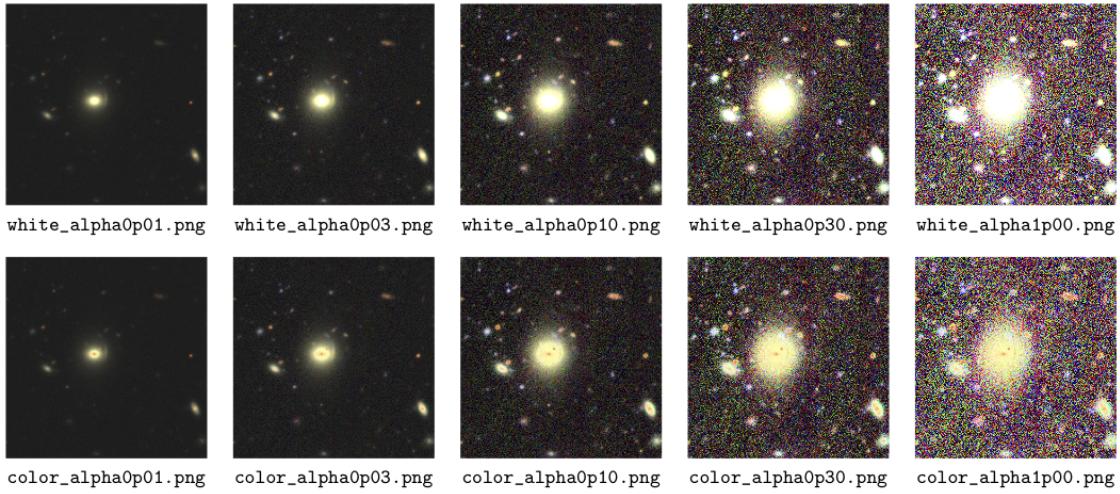


Figure 3: The effect of the saturation scheme on an example image from the CFHTLS survey. Left to right, α increases through the set $\{0.01, 0.03, 0.1, 0.3, 1.0\}$. Top row: saturation to white; Bottom row: color-preserving saturation, as proposed by L04.

At the bright end we have a different choice to make: what to do in pixels whose values are greater than 1 in any channel? In Figures 1 and 2 we simply snapped these pixel values to one in that channel, a procedure that leads to “saturation to white” in the case where all three channel pixel values exceed 1. An alternative is to snap the highest pixel value of an RGB triplet to 1, and then rescale the other two so as to preserve the *color* of the pixel. This was advocated in L04. The two approaches are shown in Figure 3, again as a function of α . If the image is not stretched too hard, saturation to white is not an issue, as all pixels remain in the required 0:1 range. Note that the faint objects in the low α frames in Figure 3 look very similar between the two saturation schemes. With color-preserving saturation we see some odd effects: red central cores and ring-like artifacts which, while providing more information about the images, may provide distractions during a search for low surface brightness features. The contrasting (eg yellow) rings around faint (eg red) objects is likely to be a result of PSF mismatch between these images: if the resolutions of the three channels’ images are not well-matched, confusing artifacts will arise. Saturation to white seems to be an *easy way to hide this problem*.

4 Color Balance

Finally, we return to the choice of scales for a composite image, which is best made after setting the stretch and saturation parameters well. The scales should reflect the quality of each channel’s image, including the sensitivity of the instrumentation, the filter transmission, the exposure time and so on. However, the relative scales can also be chosen to change the balance of color in the image, in order to improve the color contrast between different objects. In Figure 4 we show the effect of varying the scales by small amounts around their natural (unit) values.

A good strategy when looking for contrasting features around massive galaxies could be to choose

scales such that the massive galaxies appear as bright yellow as possible, and the objects around them as different as possible. The set $\{0.8, 1.0, 1.0\}$ seems to work reasonably well in this example.

5 Conclusions

From these explorations we conclude:

- α is the key “contrast” parameter needed to bring out low surface brightness features in an image.
- If all images in a set were taken under the same conditions, only the relative scales need be specified.
- These relative scales determine the color balance of the image, and should be chosen by experimentation.
- The Q parameter is a “brightness” control, and has less effect on an image than α ; however, they do need to be set together.
- Future work using images with varying conditions may need unnormalized scales, in which case Q may become somewhat, although not completely, redundant.
- While color-preserving saturation retains more information in the image, this information may be distracting during a low surface brightness feature search.
- Matching the resolutions of the input channel images may mitigate against some of the artifacts highlighted by the saturate to color algorithm.

Further work should include allowing unnormalized scales, one per image, to allow for images with different units, as concluded above. Then we should enable combinations of $N > 3$ images, breaking the one image per channel paradigm set up here. Input images will need to be scaled first, then combined into channel images, and then stretched as described in Section 2.

6 References

References

Lupton R., Blanton M. R., Fekete G., Hogg D. W., O’Mullane W., Szalay A., Wherry N., 2004, PASP, 116, 133

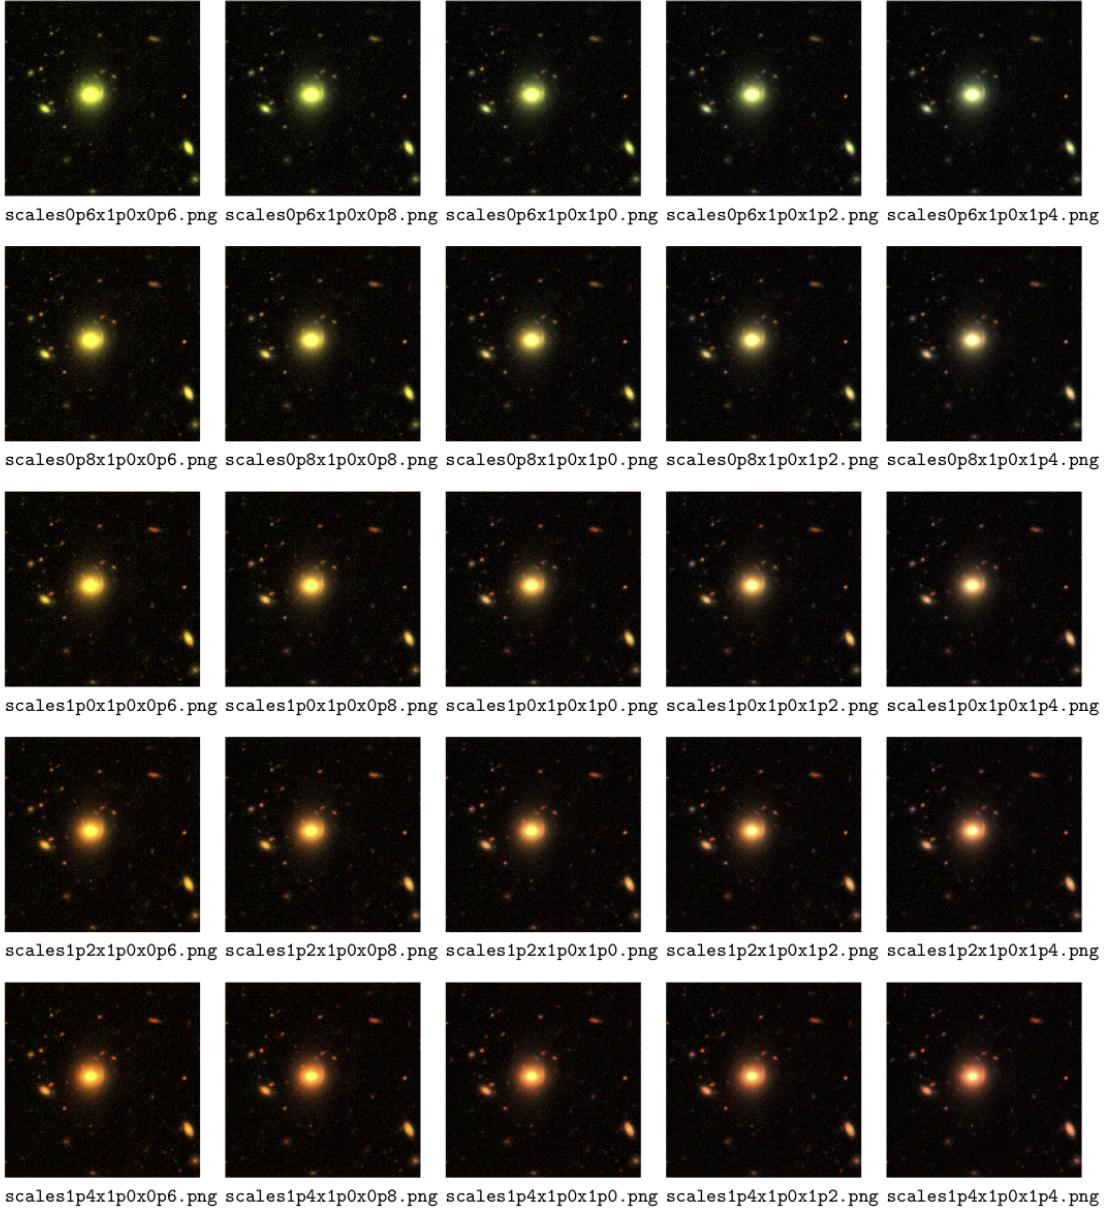


Figure 4: The effect of changing the relative image scales on the color balance in an example image from the CFHTLS survey. We keep the G channel image scale fixed at 1.0. Left to right, the R channel image scale increases through the set $\{0.6, 0.8, 1.0, 1.2, 1.4\}$. Top to bottom, the B channel image increases through the same set. $Q = 1.0$ and $\alpha = 0.03$.